

Concurrent Computations on Multicore Processors

Threads & Subprocesses

Christian Zielinski

18th June 2015 – PyCon SG 2015

Who am I?

- Currently a Ph.D. student at Nanyang Technological University
 - Working in computational physics, i.e. simulation of elementary particles with Monte Carlo methods (“lattice quantum chromodynamics”)
- Heavy user of Python, **numpy** and **scipy** for data analysis and evaluation of numerical algorithms
- First experiments with Python around ten years ago
- Slides available on:

www.github.com/czielinski/pyconsg2015

Why concurrency?

- Python is *awesome*, but slow (if you don't use C/C++ extensions)
- Single CPU core performance only improves slowly over time
 - But modern processors have increasing number of cores
- Blocking operations like I/O cause idle time
 - Rather do computations in the meantime
- Concurrency helps to mitigate all problems
 - Potentially significant speed ups

Concurrency in Python

- How to make use of concurrency in Python?
- For stock Python typically use either
 - Threads via the `threading` module
 - Subprocesses via the `multiprocessing` module
- They provide convenient high-level interfaces
 - `threading` and `multiprocessing` have similar API
- There are great third-party modules to go beyond a single CPU:
 - `mpi4py`, `pycuda`, `pyopencl` (but not today's topic)

Global Interpreter Lock (GIL)

- The Global Interpreter Lock (GIL) prevents that more than one thread is executing Python bytecode at the same time
- The Global Interpreter Lock:
 - makes the CPython interpreter explicitly thread-safe
 - simplifies CPython's code base
 - prevents true parallel threading
 - is released during I/O
- The GIL is implementation dependent
 - CPython and PyPy have a GIL
 - Jython and IronPython have no GIL
- Read more on <https://wiki.python.org/moin/GlobalInterpreterLock>

Threads

- Threads are accessible via `threading` module
 - Can be spawned quickly
 - Share memory
- CPython's threads are real POSIX/Windows threads
- Threads can be useful for I/O heavy problems
 - Do useful operations during blocking operations
- Usually not suited for CPU bound problems
 - GIL prevents true parallel computations (with a few exceptions, e.g. some `numpy` functions)
 - Introduce additional overhead
- Use a `Lock` to define critical sections, i.e. code that requires mutual exclusion of access

Simple threading example

```
1 import threading as th
2
3 def worker(lock, num):
4     cubed = num**3
5     # A worker should actually avoid I/O ...
6     with lock:
7         print("{} cubed is {}".format(num, cubed))
8
9 lock = th.Lock()
10 my_threads = []
11 for i in range(4):
12     t = th.Thread(target=worker, args=(lock, i))
13     my_threads.append(t)
14     t.start()
```

Subprocesses (I)

- Subprocesses accessible via `multiprocessing` module
 - Self-sufficient interpreter instance
 - No memory sharing
- Subprocesses have their own interpreter instance
 - Spawning takes long time, significant overhead
 - Every subprocess has own GIL, allow for true parallelism

Subprocesses (II)

- Subprocesses suited for CPU bound problems
 - Can do computations truly in parallel to utilize several cores
 - Usually not useful to spawn more than one subprocess per *physical* core
- Not suited for I/O bound problems
 - Cannot e.g. read files faster from HDD
 - To keep cores busy need to feed data quickly enough
 - Potential memory bandwidth bottleneck for large number of cores

Simple multiprocessing example

```
1 import multiprocessing as mp
2
3 def worker(lock, num):
4     cubed = num**3
5     # A worker should actually avoid I/O ...
6     with lock:
7         print("{} cubed is {}".format(num, cubed))
8
9 lock = mp.Lock()
10 my_procs = []
11 for i in range(4):
12     p = mp.Process(target=worker, args=(lock, i))
13     my_procs.append(p)
14     p.start()
```

Scaling and Amdahl's law

- Using N cores does not automatically reduce runtime by a factor of N
 - Some code section inherently serial, like I/O
 - Some algorithms are difficult to parallelize, e.g. recursive functions
 - One usually parallelizes only the performance-critical parts of the code
- Amdahl's law [Amdahl '67]
 - Parallelizing a program with serial runtime T_1 using N processes and a serial code fraction $f \in [0, 1]$ (by runtime) results in a runtime of:

$$T_N = T_1 \cdot \left(f + \frac{1}{N} (1 - f) \right)$$

- So for sufficiently large N , runtime is dominated by serial code sections!

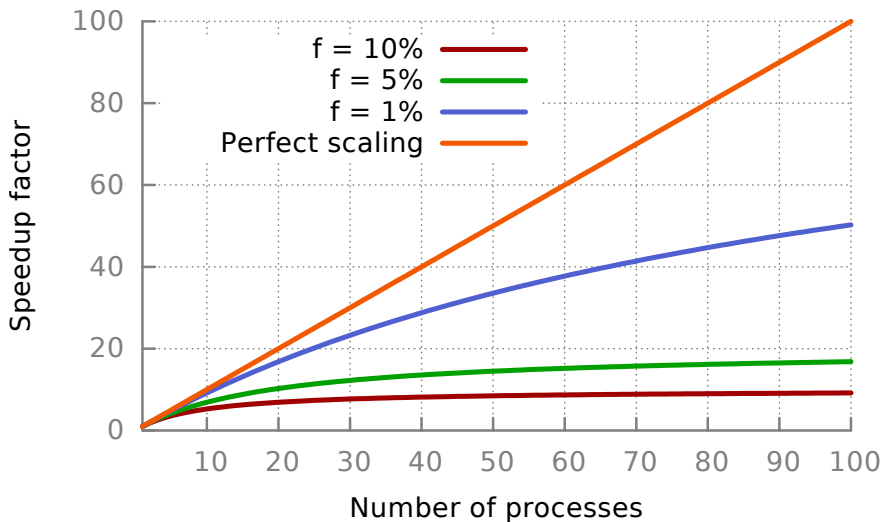


Figure: Amdahl's law

Some general remarks

- Split up larger problems into smaller subproblems with appropriate size
 - If too small, performance degraded due to overhead
 - If too large, load balancing becomes problematic
- Arguing about the correctness of a parallel program is hard
 - Order of computations are not fixed

API of the multiprocessing module

Worker pools

- Simplest parallelization strategy is to use a worker pool
 - Accessible via `multiprocessing.Pool`
- Important methods:
 - `map` (`map_async`) for (asynchronous) parallel map
 - `apply_async` for asynchronous function evaluation
- Limitations:
 - Can only deal with pickable objects
 - Cannot use `lambda` expressions, nested functions, class methods
 - Can use global functions, global classes and partial functions via `functools.partial`
 - `Pool` will not work in interactive sessions
(`__main__` module has to be importable by the children)

Parallel map

```
1 from multiprocessing import Pool
2
3 def cube(num):
4     return num**3
5
6 if __name__ == '__main__':
7     # Can also specify 'processes' parameter
8     p = Pool()
9     # Equivalent of serial map(cube, range(10))
10    res = p.map(cube, range(10))
11    print(res)
12    # Output: [0, 1, 8, 27, 64, 125, 216, 343, 512, 729]
```

Note: This overly simplified example runs slower than the serial version

Asynchronous function evaluation

```
1 from multiprocessing import Pool
2
3 def double(arg):
4     return 2*arg
5
6 def cube(num):
7     return num**3
8
9 if __name__ == '__main__':
10     p = Pool()
11     f1_get = p.apply_async(cube, args=(6,))
12     f2_get = p.apply_async(double, args=('haha',))
13     # Do other work here
14     res = [f1_get.get(), f2_get.get()]
15     print(res)
16     # Output: [216, 'hahahaha']
```

Sharing data

Shared memory maps

To share state with a `Process` use shared memory maps:

```
1 from multiprocessing import Process, Array
2
3 def cube_arr(arr):
4     for i in range(len(arr)):
5         arr[i] = arr[i]**3
6
7 if __name__ == '__main__':
8     shared_arr = Array('d', range(10)) # 'd' for double
9     p = Process(target=cube_arr, args=(shared_arr,))
10    p.start()
11    # Do other work here
12    p.join()
13    print(shared_arr[:])
14    # Output: [0, 1, 8, 27, 64, 125, 216, 343, 512, 729]
```

Server processes and managers

Alternatively use a **Manager** for more complex objects:

```
1 from multiprocessing import Process, Manager
2
3 def modify_names(names):
4     names['Alice'] = 42
5     if 'Bob' in names:
6         shared_dict['Bob'] += 1
7
8 if __name__ == '__main__':
9     with Manager() as m:
10         shared_dict = m.dict()
11         shared_dict['Bob'] = 7
12         p = Process(target=modify_names, args=(shared_dict,))
13         p.start()
14         # Do other work here
15         p.join()
16         print(shared_dict) # Output: {'Bob': 8, 'Alice': 42}
```

Interprocess communication

Pipes

To communicate between two processes we use a Pipe:

```
1 from multiprocessing import Process, Pipe
2
3 def cube(my_pipe):
4     data = my_pipe.recv()
5     my_pipe.send([n**3 for n in data])
6
7 if __name__ == '__main__':
8     parent_pipe, child_pipe = Pipe()
9     p = Process(target=cube, args=(child_pipe,))
10    p.start()
11    parent_pipe.send(range(10))    # Send work
12    # Do other work here
13    print(parent_pipe.recv())
14    p.join()
15    # Output: [0, 1, 8, 27, 64, 125, 216, 343, 512, 729]
```

Queues

For several processes we can use a Queue:

```
1 from multiprocessing import Process, Queue
2
3 def place_work(q):
4     q.put(range(10))
5
6 def process_work(q):
7     data = q.get()
8     q.put([n**3 for n in data])
9
10 if __name__ == '__main__':
11     queue = Queue()
12     p1 = Process(target=place_work, args=(queue,))
13     p2 = Process(target=process_work, args=(queue,))
14     p1.start(); p2.start(); p1.join(); p2.join()
15     print(queue.get())
16     # Output: [0, 1, 8, 27, 64, 125, 216, 343, 512, 729]
```

Guidelines (I)

- Most important: Only parallelize where necessary!
- Try to avoid sharing state if possible
- If need to share:
 - Use **Value** and **Array** for simple data
 - Use server process **Manager** for more complex objects
- For interprocess communication use **Pipe** and **Queue**
 - Pipes for fast point-to-point communication
 - Queues are multi-producer, multi-consumer FIFO queues

Guidelines (II)

- Explicitly pass resources to subprocesses and avoid global variables as they can lead to inconsistencies (Windows)
- Ensure that the main module can be safely imported by a new Python instance (Windows)
 - Use `if __name__ == '__main__':` guard to avoid side effects
- Avoid `Process.terminate` as it might break access to shared resources
- Read more on:
<https://docs.python.org/2/library/multiprocessing.html#programming-guidelines>

Time for a short demo!

- Midpoint rule for numerical integration

$$\int_a^b f(x) \, dx \approx h \sum_{k=1}^N f\left(a + \left(k - \frac{1}{2}\right) h\right), \quad h = \frac{b-a}{N}$$

- Allows for straightforward parallel evaluation

$$\int_a^b f(x) \, dx = \int_a^c f(x) \, dx + \int_c^b f(x) \, dx, \quad \forall c \in \mathbb{R}$$

- Can split up large integration range into several small ones, which can be evaluated in parallel

Questions & Answers

(if time permits)

Slides on:

`www.github.com/czielinski/pyconsg2015`